

Hi everyone.

This notebook is part of a project developed for the Landing.Jobs Data Challenge

If you want to read more about it, you can read the abstract and the presentation through these links:

Abstract: https://drive.google.com/file/d/1qeOC-0TPkGPZHNIgGlaJ6g-cPKNCuAP_/view?usp=sharing
(https://drive.google.com/file/d/1qeOC-0TPkGPZHNIgGlaJ6g-cPKNCuAP_/view?usp=sharing).

Presentation: https://drive.google.com/file/d/1eCMqhjT_xGBNw_JSAkUrW_FgkNSMtWlv/view?usp=sharing
(https://drive.google.com/file/d/1eCMqhjT_xGBNw_JSAkUrW_FgkNSMtWlv/view?usp=sharing).

Goal: Developing a model able to predict if a tech professional earns more than 30k€. This model aims to be ethical and will help anyone that wants to ask for a raise, regardless the gender, nationality and age.

To start, let's import all the libraries that will help us to build the model and do data transformation

```
In [1]:  
  
import pandas as pd  
import numpy as np  
from numpy import mean  
from numpy import isnan  
from numpy import asarray  
from numpy import polyfit  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.preprocessing import OneHotEncoder  
from sklearn.preprocessing import MinMaxScaler  
from sklearn.linear_model import LogisticRegression  
from numpy import std  
from sklearn.datasets import make_classification  
from sklearn.model_selection import cross_val_score  
from sklearn.model_selection import RepeatedStratifiedKFold  
from sklearn.model_selection import train_test_split  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import confusion_matrix  
from sklearn import metrics  
from sklearn.model_selection import KFold  
from sklearn.model_selection import LeaveOneOut  
from sklearn.linear_model import RidgeClassifier  
from sklearn.linear_model import SGDClassifier  
from sklearn.linear_model import PassiveAggressiveClassifier  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.tree import ExtraTreeClassifier  
from sklearn.svm import LinearSVC  
from sklearn.svm import SVC  
from sklearn.naive_bayes import GaussianNB  
from sklearn.ensemble import AdaBoostClassifier  
from sklearn.ensemble import BaggingClassifier  
from sklearn.ensemble import ExtraTreesClassifier  
from sklearn.gaussian_process import GaussianProcessClassifier  
from sklearn.ensemble import GradientBoostingClassifier  
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis  
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis  
from sklearn.model_selection import GridSearchCV  
from xgboost import XGBClassifier
```

Brief Data Analysis

- I really recommend to read the Lanfing.Jobs Report 2021. A lot of decisions were made based on the report analysis.

We need to understand the data we are working with. In this section I will delete irrelevant columns to solve this problem, due to not presenting any value or because they may have duplicated information in some way.

```
In [2]:
df = pd.read_csv('LJDV.csv', sep = ';', )
df.index = df.ID
df.shape
```

```
(3371, 126)
```

finding columns that might be irrelevant, having only one unique value

```
In [3]:
# All Language and Framework columns are expected to have only 1.
#But some have 0 and I'll drop those as well
df.nunique()[df.nunique() < 2]
```

```
Residence_Country          1
Language_JavaScript         1
Language_Bash/Shell/PowerShell  1
Language_SQL                1
Language_Java               1
Language_C#                 1
Language_Python             1
Language_PHP                1
Language_C++                1
Language_C                  1
Language_TypeScript         1
Language_Ruby               1
Language_Swift              1
Language_Objective-C        1
Language_VB.NET             1
Language_Assembly           1
Language_R                  1
Language_Perl               1
Language_VBA                1
Language_Matlab             1
Language_Go                 1
Language_Scala              1
Language_Groovy             1
Language_Coffee Script     1
Language_Visual Basic 6    1
Language_Lua                1
Language_Haskell            1
Language_HTML/CSS          1
Language_Kotlin             1
Language_Rust               1
Language_Elixir             1
Language_Clojure            1
Language_WebAssembly        1
Language_Dart               1
Language_Languages_N/A      0
Framework_jQuery            1
Framework_.NET              1
Framework_Angular/Angular.js  1
Framework_Ruby on Rails     1
Framework_React             1
Framework_Django            1
Framework_Laravel           1
Framework_Spring            1
Framework_Vue.js            1
Framework_Express           1
Framework_Meteor             1
Framework_Flask              1
Framework_EMBER.js          1
Framework_Drupal            1
Framework_OutSystems        1
Framework_Framework_N/A     0
dtype: int64
```

```
In [4]:
df = df.drop(['ID', 'Work_as_Contractor_12m', 'Work_as_Perm', 'Remote_Working_Current', 'Remote_Working_Current_Flexible_Office_Days', 'Remote_Working_due_to_Covid', 'Remote_Work_Opinion', 'Residence_Country', 'Language_Languages_N/A', 'Framework_Framework_N/A', 'Employment_Status_Aggregated', 'Residence_District_Aggregated', 'Work_Company_Country', 'Work_Company_Continent', 'Work_Company_PT_District_Aggregated', 'Remote_Working_Current_Flexible_Office_Days', 'Job_Role_Original', 'Job_Role_Other', 'Employer_Industry_Other', 'Employer_Org_Type_Other', 'Contractor_Avg_Project', 'Contractor_Avg_Annual_Salary', 'Perm_GAS_Avg', 'Perm_GAS_Low_Limit', 'Perm_GAS_High_Limit', 'Avg_Salary', 'Salary_Change'], axis=1)
df = df.drop(['Salary_Fairness', 'Changing_Jobs_next_6_months', 'Job_Motivator_Work_life_balance', 'Job_Motivator_Compensation_and_benefits', 'Job_Motivator_Training/Development_programs_at_work', 'Job_Motivator_Career_growth_opportunities', 'Job_Motivator_Remote_working', 'Job_Motivator_Flexible_schedule', 'Job_Motivator_Company_culture', 'Job_Motivator_The_technologies_I'm_working_with', 'Job_Motivator_Versatility/Variety_of_projects', 'Job_Motivator_Freedom_to_choose_the_clients_and/or_projects', 'Job_Motivator_Being_autonomous_at_work', 'Job_Motivator_How_widely_used_or_impactful_the_product/service_I_work_on_is', 'Job_Motivator_Environmentally_friendly/responsible_work_practice', 'Job_Perk_Meals_allowance/Company_provided_meals_or_snacks', 'Job_Perk_Transportation_benefit', 'Job_Perk_Health_benefits', 'Job_Perk_Fitness_or_wellness_benefit_(ex._gym_membership)', 'Job_Perk_Computer/Office_equipment_allowance', 'Job_Perk_Professional_development_sponsorship', 'Job_Perk_Annual_bonus', 'Job_Perk_Long-term_leave', 'Job_Perk_Parental_leave', 'Job_Perk_Stock_options_or_shares', 'Job_Perk_Education_sponsorship', 'Job_Perk_Child_care'], axis=1)
df = df.drop(['Birth_Year', 'Way_Into_Tech_Other', 'Working_Experience_Aggregated'], axis=1)
df = df.drop(['Language_Languages_Other', 'Framework_Framework_Other'], axis=1)
df = df.drop(['Contractor_Avg_Project_Intervals', 'Contractor_Avg_Hour_Rate'], axis=1)
```

From now on, I'll work with a dataset only with Full time employees

```
In [5]:
df_ft = df[df['Employment_Status'] == 'Employed full-time'].drop(['Employment_Status'], axis = 1)
```

```
In [6]:
df_ft.nunique()[df_ft.nunique() < 1]
```

```
Series([], dtype: int64)
```

Data Transformation

Due to the nature of the features, we have a lot of nan values. For the model to process the data, I'll convert all nan values to 0 or even remove some rows, that can create samples with incomplete information. The Target variable will be the "Per_GAS". Below 30k will be 0, Above 30k will be 1. In this way, I create a binary classification problem.

Our dataset has 2931 rows right now. The Target variable will be divided with 1669 samples that are below 30k, target '0', and 1262 samples that are above 30k, target '1'.

- 1: 43.06%
- 0: 56.94%

Categorical Columns - Some columns have a hierarchical value and for those cases I'll use numeric labeling to keep their value. In other cases, those columns are purely categorical and the solution will be one hot encoding

In the end will also change the data type of all the dataset to float and normalize the data

A little disclaimer: Haven't removed any outlier, because removing them will harm the model. I've tried the isolation forest and z-score methods.

```
In [7]:
df_ft['Work_Company_PT_District'].replace(np.nan, 0, inplace=True)
```

```
In [8]:
```

```
df_ft.isnull().sum()[df_ft.isnull().sum() > 0]
```

Employer_Industry	102
Employer_Org_Type	30
Language_JavaScript	1541
Language_Bash/Shell/PowerShell	2564
Language_SQL	1670
Language_Java	2377
Language_C#	2428
Language_Python	2565
Language_PHP	2708
Language_C++	2976
Language_C	3040
Language_TypeScript	2473
Language_Ruby	2988
Language_Swift	3028
Language_Objective-C	3080
Language_VB.NET	3027
Language_Assembly	3120
Language_R	3072
Language_Perl	3104
Language_VBA	3038
Language_Matlab	3111
Language_Go	3011
Language_Scala	3065
Language_Groovy	3057
Language_Coffee Script	3122
Language_Visual Basic 6	3089
Language_Lua	3114
Language_Haskell	3129
Language_HTML/CSS	2014
Language_Kotlin	2961
Language_Rust	3109
Language_Elixir	3100
Language_Clojure	3124
Language_WebAssembly	3116
Language_Dart	3090
Framework_jQuery	2410
Framework_.NET	2404
Framework_Angular/Angular.js	2598
Framework_Ruby on Rails	2998
Framework_React	2447
Framework_Django	3004
Framework_Laravel	2954
Framework_Spring	2745
Framework_Vue.js	2870
Framework_Express	2898
Framework_Meteor	3126
Framework_Flask	3031
Framework_EMBER.js	3115
Framework_Drupal	3088
Framework_OutSystems	3083
Age	4
Way_Into_Tech	69
Education_Level	2

dtype: int64

```
In [9]:
```

```
df_ft = df_ft[df_ft['Employer_Industry'].notna()]
df_ft = df_ft[df_ft['Employer_Org_Type'].notna()]
df_ft = df_ft[df_ft['Age'].notna()]
df_ft = df_ft[df_ft['Way_Into_Tech'].notna()]
df_ft = df_ft[df_ft['Education_Level'].notna()]
```

```
In [10]:
columns_0 = df_ft.isnull().sum()[df_ft.isnull().sum() > 0].index
df_ft = df_ft.fillna(0)
df_help = df_ft.copy()
df_ft
```

	Residence_District	Work_Company_PT_International	Work_Company_PT_District	Job_Remote_or_Office	Job_Role
ID					
4	Açores	Portugal	Açores	Full Office Job	Product Owner/Product Manager
5	Braga	Portugal	Braga	Remote Job (full or flexible)	Back-End Developer
7	Aveiro	Portugal	Aveiro	Full Office Job	Front-End Developer
8	Santarém	Portugal	Lisboa	Full Office Job	Full-Stack Developer
9	Viseu	Portugal	Viseu	Full Office Job	Full-Stack Developer
...
3365	Lisboa	International	0	Remote Job (full or flexible)	Maintenance & Support
3367	Porto	Portugal	Porto	Remote Job (full or flexible)	Full-Stack Developer
3368	Lisboa	Portugal	Lisboa	Full Office Job	Quality Assurance/Testing
3369	Porto	International	0	Full Office Job	Front-End Developer
3370	Porto	Portugal	Porto	Remote Job (full or flexible)	Front-End Developer

2931 rows × 65 columns

Transforming Perm_GAS as Target with numeric values

In [11]:

```
def f(row):
    if row['Perm_GAS'] == '< €15.000':
        val = 0
    if row['Perm_GAS'] == '€15.000 - €20.000':
        val = 0
    if row['Perm_GAS'] == '€20.000 - €25.000':
        val = 0
    if row['Perm_GAS'] == '€25.000 - €30.000':
        val = 0
    if row['Perm_GAS'] == '€30.000 - €35.000':
        val = 1
    if row['Perm_GAS'] == '€35.000 - €40.000':
        val = 1
    if row['Perm_GAS'] == '€40.000 - €45.000':
        val = 1
    if row['Perm_GAS'] == '€45.000 - €50.000':
        val = 1
    if row['Perm_GAS'] == '€50.000 - €55.000':
        val = 1
    if row['Perm_GAS'] == '€55.000 - €60.000':
        val = 1
    if row['Perm_GAS'] == '€60.000 - €65.000':
        val = 1
    if row['Perm_GAS'] == '€65.000 - €70.000':
        val = 1
    if row['Perm_GAS'] == '€70.000 - €75.000':
        val = 1
    if row['Perm_GAS'] == '€75.000 - €80.000':
        val = 1
    if row['Perm_GAS'] == '€80.000 - €85.000':
        val = 1
    if row['Perm_GAS'] == '€85.000 - €90.000':
        val = 1
    if row['Perm_GAS'] == '€90.000 - €95.000':
        val = 1
    if row['Perm_GAS'] == '€95.000 - €100.000':
        val = 1
    if row['Perm_GAS'] == '> €100.000':
        val = 1
    return val
```

In [12]:

```
df_ft['Target'] = df_ft.apply(f, axis=1)
df_ft = df_ft.drop(['Perm_GAS'], axis = 1)
df_ft['Target'].value_counts().sort_values()
```

```
1    1262
0    1669
Name: Target, dtype: int64
```

In [13]:

```
df_ft.shape
```

```
(2931, 65)
```

In [14]:

```
df_ft.Target.isnull().sum()
```

```
0
```

Solving categorical columns Problem

creating array with categorical columns

```
In [15]:  
cat_col = ['Residence_District', 'Work_Company_PT_District', 'Work_Company_PT_International', 'Job_Remote_or_Office', 'Job_Role', 'Employer_Industry', 'Employer_Org_Type', 'Employer_Size', 'Way_Into_Tech', 'Education_Level', 'Perm_Current_Company_how_long', 'Citizenship', 'Gender', 'English_Level', 'Working_Experience']
```

Creating Dummies for each categorical column

```
In [16]:  
  
#Creating new columns for each element  
encoded_residence = pd.get_dummies(df_ft['Residence_District'], prefix = 'Colab_Resid')  
#joining the results to our dataset  
df_ft = df_ft.join(encoded_residence).drop('Residence_District', axis=1)
```

```
In [17]:  
  
#Creating new columns for each element  
Work_Company_PT_District = pd.get_dummies(df_ft['Work_Company_PT_District'], prefix = 'Comp_District')  
#joining the results to our dataset  
df_ft = df_ft.join(Work_Company_PT_District).drop('Work_Company_PT_District', axis=1)
```

```
In [18]:  
  
#Creating new columns for each element  
Work_Company_PT_International = pd.get_dummies(df_ft['Work_Company_PT_International'], prefix = 'Comp_Country')  
#joining the results to our dataset  
df_ft = df_ft.join(Work_Company_PT_International).drop('Work_Company_PT_International', axis=1)
```

```
In [19]:  
  
#Creating new columns for each element  
Job_Remote_or_Office = pd.get_dummies(df_ft['Job_Remote_or_Office'], prefix = 'Job_rem_off')  
#joining the results to our dataset  
df_ft = df_ft.join(Job_Remote_or_Office).drop('Job_Remote_or_Office', axis=1)
```

```
In [20]:  
  
#Creating new columns for each element  
Job_Role = pd.get_dummies(df_ft['Job_Role'], prefix = 'job_role')  
#joining the results to our dataset  
df_ft = df_ft.join(Job_Role).drop('Job_Role', axis=1)
```

```
In [21]:  
  
#Creating new columns for each element  
Employer_Industry = pd.get_dummies(df_ft['Employer_Industry'], prefix = 'Employer_Industry')  
#joining the results to our dataset  
df_ft = df_ft.join(Employer_Industry).drop('Employer_Industry', axis=1)
```

```
In [22]:  
  
#Creating new columns for each element  
  
Employer_Org_Type = pd.get_dummies(df_ft['Employer_Org_Type'], prefix = 'Type')  
  
#there is a column with a character "<", this character creates an error in some models. Will re  
move it:  
Employer_Org_Type['Type_SME - Small or Medium Enterprise (personnel <250)'] = Employer_Org_Type[  
'Type_SME - Small or Medium Enterprise (personnel <250)']  
Employer_Org_Type = Employer_Org_Type.drop(['Type_SME - Small or Medium Enterprise (personnel <2  
50)'], axis = 1)  
  
#joining the results to our dataset  
df_ft = df_ft.join(Employer_Org_Type).drop('Employer_Org_Type', axis=1)
```

```
In [23]:  
  
#In this case, for "Employer Size", because we can translate this variable into numerical value,  
I will attribute a label for each size level  
  
df_ft['Employer_Size'] = df_ft['Employer_Size'].map({'Less than 10 employees': 0, '10 - 19 emplo  
yees': 1,  
                                                    '20 - 99 employees': 2, '100 - 499 employees': 3,  
                                                    '500 - 999 employees': 4, '1000 - 4.999 employees': 5, 'More  
than 5.000 employees': 6})
```

```
In [24]:  
  
#Creating new columns for each element  
Way_Into_Tech = pd.get_dummies(df_ft['Way_Into_Tech'], prefix = 'Way_Into_Tech')  
#joining the results to our dataset  
df_ft = df_ft.join(Way_Into_Tech).drop('Way_Into_Tech', axis=1)
```

```
In [25]:  
  
#In this case, for "Education_Level", because we can translate this variable into numerical valu  
e, I will attribute a label for each size level  
  
df_ft['Education_Level'] = df_ft['Education_Level'].map({'I prefer not to answer': 0, 'Basic Edu  
cation': 1,  
                                                         'High School Education': 2, 'Trade/technical/vocational trai  
ning': 3,  
                                                         'University drop out': 4, 'Bachelor degree': 5, 'Masters deg  
ree': 6, 'Doctoral degree': 7})
```

```
In [26]:  
  
#In this case, for "Perm_Current_Company_how_long", because we can translate this variable into  
numerical value, I will attribute a label for each size level  
  
df_ft['Perm_Current_Company_how_long'] = df_ft['Perm_Current_Company_how_long'].map({'Less than  
one year': 0, 'Between 1 - 3 years': 1,  
                                             'More than 3 years': 2})
```

```
In [27]:  
  
#Beacause of the ethics associated with this model, this feature will be removed.  
df_ft = df_ft.drop('Citizenship', axis=1)
```

```
In [28]:  
  
#Beacause of the ethics associated with this model, this feature will be removed.  
df_ft = df_ft.drop('Gender', axis=1)
```


In [29]:

```
#In this case, for "English_Level", because we can translate this variable into numerical value,
I will attribute a label for each size level
```

```
df_ft['English_Level'] = df_ft['English_Level'].map({'Elementary': 0, 'Limited working proficiency': 1,
                                                    'Professional working proficiency': 2, 'Full professional proficiency': 3, 'Native or bilingual proficiency': 4})
```

In [30]:

```
#In this case, for "Working_Experience", because we can translate this variable into numerical value,
I will attribute a label for each size level
```

```
df_ft['Working_Experience'] = df_ft['Working_Experience'].map({'No working experience': 0, 'Less than 1 year': 1,
                                                              'Between 1 - 3 years': 2, 'Between 3 - 6 years': 3, 'Between 6 - 9 years': 4, 'More than 9 years': 5})
```

In [31]:

```
#Beacause of the ethics associated with this model, this feature will be removed.
```

```
df_ft = df_ft.drop('Age', axis=1)
```

In [32]:

```
#defining all string as 1 in languages and frameworks columns
```

```
#df_help1 = df_help.iloc[:, 8:56]
```

```
df_ft['Employer_Size_'] = df_ft['Employer_Size'].copy()
```

```
df_ft = df_ft.drop(['Employer_Size'], axis = 1)
```

```
#creating a new dataset with language and framework to turn them to numeric
```

```
df_help1 = df_ft.iloc[:, 0:48]
```

```
#New column with the sum of how many languages and frameworks does one knows
```

```
df_help1[df_help1 != 0] = 1
```

```
df_ft = df_ft.iloc[:, 48:154]
```

```
#join both datasets again into one
```

```
df_ft = df_ft.join(df_help1)
```

```
df_ft.shape
```

```
(2931, 152)
```

Changing data Types

In [33]:

```
df_ft = df_ft.astype(float)
```

Dropping Unnecessary columns

In [34]:

```
#when the company is international
```

```
df_ft['Comp_District_0'].sum()
```

```
#repeated informatino in columns that tell us if it is remote or not
```

```
df_ft = df_ft.drop(['Comp_District_0', 'Comp_Country_Portugal', 'Job_rem_off_Remote Job (full or flexible)'], axis = 1)
```

```
In [35]:
df_ft
```

	Perm_Current_Company_how_long	English_Level	Education_Level	Working_Experience	Target	Colab_Resid_Aveiro	Ci
ID							
4	1.0	3.0	5.0	5.0	0.0	0.0	1.0
5	2.0	2.0	5.0	4.0	0.0	0.0	0.0
7	1.0	3.0	5.0	3.0	0.0	1.0	0.0
8	2.0	4.0	5.0	5.0	0.0	0.0	0.0
9	1.0	1.0	5.0	2.0	0.0	0.0	0.0
...
3365	1.0	3.0	4.0	5.0	1.0	0.0	0.0
3367	1.0	3.0	4.0	3.0	0.0	0.0	0.0
3368	1.0	2.0	4.0	3.0	0.0	0.0	0.0
3369	2.0	2.0	4.0	4.0	1.0	0.0	0.0
3370	1.0	3.0	4.0	3.0	0.0	0.0	0.0

2931 rows × 149 columns

Normalizing Data

```
In [36]:
scaler = MinMaxScaler()
#creating new variable with target column to keep as discrete and not continuous
y = df_ft['Target'].astype(int)
df_ft = df_ft.drop(['Target'], axis = 1)
df_ft = pd.DataFrame(scaler.fit_transform(df_ft), columns=df_ft.columns)
df_ft['Target'] = y.reset_index().drop(['ID'], axis = 1).Target
y = y.reset_index().drop(['ID'], axis = 1).Target
```

```
In [37]:
X = df_ft.drop(['Target'], axis = 1)
```

Creating Training and Test set

For this model, will be create a Train and validation test. The final score (accuracy) will be validated with cross validation.

```
In [38]:
#Creating test and train dataset
x_train, x_val, y_train, y_val = train_test_split(X, y, test_size=0.20, random_state=0)

x_train = x_train.reset_index().drop(['index'], axis = 1)
x_val = x_val.reset_index().drop(['index'], axis = 1)

y_train = y_train.reset_index().Target
y_val = y_val.reset_index().Target
```

```
In [39]:
x_train.shape

(2344, 148)
```

```
In [40]:
x_val.shape

(587, 148)
```

What model to use

In order to know what model to use and to start to understand the data to extract the best accuracy possible, I'll run several models to know how will they perform. To choose the final model, I'll always have in consideration that I want not just the best model, but the best model with a good level of interpretability from their coefficients.

get a list of models to evaluate

```
def get_models(): models = list() models.append(LogisticRegression()) models.append(RidgeClassifier())
models.append(SGDClassifier()) models.append(PassiveAggressiveClassifier()) models.append(KNeighborsClassifier())
models.append(DecisionTreeClassifier()) models.append(ExtraTreeClassifier()) models.append(LinearSVC())
models.append(SVC()) models.append(GaussianNB()) models.append(AdaBoostClassifier())
models.append(BaggingClassifier()) models.append(RandomForestClassifier()) models.append(ExtraTreesClassifier())
models.append(GaussianProcessClassifier()) models.append(GradientBoostingClassifier())
models.append(LinearDiscriminantAnalysis()) models.append(QuadraticDiscriminantAnalysis()) return models

def evaluate_model(cv, model): scores = cross_val_score(model, x_train, y_train, scoring='accuracy', cv=cv, n_jobs=-1)
return mean(scores)

cv_2 = KFold(n_splits=5, shuffle=True, random_state=1) cv = KFold(n_splits=20, shuffle=True, random_state=1) models =
get_models() cv_2_results, cv_results = list(), list() for model in models: cv_mean = evaluate_model(cv, model) cv_2_mean
= evaluate_model(cv_2, model) if isnan(cv_mean) or isnan(cv_2_mean): continue cv_results.append(cv_mean)
cv_2_results.append(cv_2_mean) print('> %s: cv_2=%0.3f, cv=%0.3f' % (type(model).name, cv_2_mean, cv_mean))
```

In []:

Most important features / Feature reduction

I'll try to understand what are the right features to integrate in the model. For this case, I'll also bear in mind that PCA is a technic that can possibly improve accuracy but will lack interpretability. The final results will show that PCA results do in fact improve the model, but not so much. For that reason, I'll choose Logistic Regression and the features included will be the number of features of higher coefficient that return the highest accuracy rate.

Top features from PCA on Logistic Regression

In [41]:

```
#I'll try to understand what are the best fetures to optimize the model from PCA
from sklearn.decomposition import PCA

score_list = []
for x in np.arange(1, x_train.shape[1]):
    #Defining the columns set
    pca = PCA(n_components = x)
    pca.fit(x_train)
    x_trainPCA = pca.transform(x_train)
    x_valPCA = pca.transform(x_val)
    #Creating Train and Test
    #Score from LogReg
    logisticRegr = LogisticRegression()
    logisticRegr.fit(x_trainPCA, y_train)
    score = logisticRegr.score(x_valPCA, y_val)
    score_list.append([score, x])
```

```
In [42]:
topPCA = pd.DataFrame(score_list).sort_values(0,ascending = False)
topPCA
```

	0	1
71	0.814310	72
72	0.809199	73
73	0.807496	74
75	0.807496	76
68	0.807496	69
...
4	0.618399	5
3	0.618399	4
2	0.592845	3
0	0.591141	1
1	0.589438	2

147 rows × 2 columns

```
In [ ]:
```

Top features from a LogReg with top correlation matrix

Building the LogReg with all columns and building a dataframe with top coefficients:

```
In [43]:
#Getting top correlated columns with the Target. As expected, because of our coefficients in Log
Reg, the highest value is for "Working Experience" with 0.48.
cor = df_ft.corr()
col3 = cor.sort_values('Target')['Target'].sort_values(ascending = False)
```

```
In [44]:
#I'll try to understand what is the optimal number of top features from the correlation matrix
(pearson) to get the best accuracy rate.

score_list = []
for x in range(2,149):
    #Defining the columns set
    cols_3 = X[col3.index[1:x]].columns
    x_train_LRcor = x_train[cols_3]
    x_val_LRcor = x_val[cols_3]
    #Creating Train and Test
    #Score from LogReg
    logisticRegr = LogisticRegression()
    logisticRegr.fit(x_train_LRcor, y_train)
    score = logisticRegr.score(x_val_LRcor, y_val)
    score_list.append(score)
```

```
In [45]:
pd.options.display.max_rows = 50
pd.DataFrame(score_list).sort_values(0, ascending = False).head(5)
```

	0
146	0.805792
122	0.800681
124	0.798978
123	0.798978
79	0.797274

Top features from a LogReg with top coef from LogReg with all columns

```
In [46]:
logisticRegr = LogisticRegression(solver='newton-cg')
logisticRegr.fit(x_train, y_train)
score = logisticRegr.score(x_val, y_val)
print(score)

pd.options.display.max_rows = 500

importance = logisticRegr.coef_

cols_importance = pd.DataFrame(importance)
imp = cols_importance.T
imp['columns'] = X.columns
imp = imp.sort_values(by = 0, ascending = False).reset_index().drop(['index'], axis = 1)

0.8006814310051107
```

```
In [47]:
#I'll try to understand what is the optimal number of top features from the logreg coefficients to get the best accuracy rate.

score_list_LR = []

for x in range(1,160):
    #Defining the columns set
    cols_2 = list(imp['columns'][0:x])
    x_train_LR = x_train[cols_2]
    x_val_LR = x_val[cols_2]
    #Score from LogReg
    logisticRegr = LogisticRegression(solver='newton-cg')
    logisticRegr.fit(x_train_LR, y_train)
    score = logisticRegr.score(x_val_LR, y_val)
    score_list_LR.append([score,x])
```

```
In [48]:
pd.options.display.max_rows = 50
top_LR = pd.DataFrame(score_list_LR).sort_values(0, ascending = False).head(5)
top_LR
```

	0	1
139	0.807496	140
140	0.807496	141
141	0.805792	142
144	0.804089	145
143	0.804089	144

Creating Our Model from previous experiences

I've also tried a model with GBoost and XGBoost that would give great results sometimes, but due to their stochastic nature, those good results were very uncertain. When I tried to stabilize the stochasticity of those models, the accuracy fell to levels below Logistic Regression.

The final model will be made of a Logistic Regression for prediction, using top features from logistic regression coefficients. It's simple and very easy to use and I hope I can optimize it a little bit.

Filtering our Train and Test sets with the best results obtained before

```
In [49]:  
cols_2 = list(imp['columns'][0:top_LR[1].iloc[0]])  
x_train = x_train[cols_2]  
x_val = x_val[cols_2]
```

```
In [50]:  
x_train.shape  
  
(2344, 140)
```

Applying the model with the new sets, should give us the best result obtained previously

```
In [51]:  
logisticRegr = LogisticRegression(solver='newton-cg')  
logisticRegr.fit(x_train, y_train)  
score = logisticRegr.score(x_val, y_val)  
score  
  
0.807495741056218
```

This step is very important and very interesting. We are now able to read the coefficients for each feature and understand their importance. Working Experience is clearly the king in our model and completely makes sense: The higher your experience, the higher will be your wage.

```
In [52]:  
  
#Finding the coefficient values for each column  
  
importance = logisticRegr.coef_  
  
cols_importance = pd.DataFrame(importance)  
imp = cols_importance.T  
imp['columns'] = x_train.columns  
imp['absolute'] = imp[0].abs()  
imp = imp.sort_values(by = 'absolute', ascending = False).reset_index().drop(['index'], axis = 1  
)  
imp.head(20)
```

	0	columns	absolute
0	4.915877	Working_Experience	4.915877
1	1.866013	Comp_Country_International	1.866013
2	1.516652	Type_Scale-up (fast growing company aka "unico...	1.516652
3	1.389163	Colab_Resid_Lisboa	1.389163
4	1.348296	Type_Startup (new business venture)	1.348296
5	1.193807	Education_Level	1.193807
6	1.055859	Comp_District_Lisboa	1.055859
7	1.031700	Employer_Size_	1.031700
8	1.024210	Colab_Resid_Setúbal	1.024210
9	1.009382	Employer_Industry_Media, advertising, publishi...	1.009382
10	0.961249	English_Level	0.961249
11	0.950513	job_role_Technical Team Leader	0.950513
12	0.942491	job_role_CTO	0.942491
13	0.910483	Comp_District_Coimbra	0.910483
14	0.809730	Employer_Industry_Transportation	0.809730
15	0.804941	Comp_District_Porto	0.804941
16	0.785593	Employer_Industry_Energy or utilities	0.785593
17	0.774407	job_role_Scrum Master	0.774407
18	0.774021	Language_Perl	0.774021
19	0.753584	Type_Corporate	0.753584

Bagging

From this point, I've tried several ways to improve the model and this one turned out to be a fast way of improving the accuracy a little bit

I've decided to use the top 3 models from the previous results from the process where I've tried several models and create a separated dataset with 3 columns, each one for the prediction of each model. The models used are:

- Logistic Regression
- XGBoost
- SVC

Creating new column with prediction probability #1

```
In [53]:  
  
x_train_pred = logisticRegr.predict_proba(x_train)[: ,1]  
x_val_pred = logisticRegr.predict_proba(x_val)[: ,1]  
  
#creating 1st prediction column  
  
train_pred = pd.DataFrame()  
val_pred = pd.DataFrame()  
test_pred = pd.DataFrame()  
  
train_pred['LR'] = x_train_pred  
val_pred['LR'] = x_val_pred
```

Creating new column with prediction probabilty #2 - using XGB

```
In [54]:  
  
from numpy import loadtxt  
from xgboost import XGBClassifier  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score  
  
model = XGBClassifier()  
model.fit(x_train, y_train)  
  
# make predictions for test data  
x_train_pred = model.predict_proba(x_train)[: ,1]  
x_val_pred = model.predict_proba(x_val)[: ,1]  
  
#creating 2nd prediction column  
  
train_pred['XGB'] = x_train_pred  
val_pred['XGB'] = x_val_pred
```

Creating new column with prediction probabilty #3 - using SVC

```
In [55]:  
  
from sklearn.svm import SVC  
from sklearn.neighbors import KNeighborsClassifier  
  
SVC2 = SVC(probability=True)  
  
SVC2.fit(x_train,y_train)  
  
# make predictions for test data  
x_train_pred = SVC2.predict_proba(x_train)[: ,1]  
x_val_pred = SVC2.predict_proba(x_val)[: ,1]  
  
#creating 3rd prediction column  
  
train_pred['SVC'] = x_train_pred  
val_pred['SVC'] = x_val_pred
```

```
In [ ]:
```

Final Model

The final model will be a Logistic regression, but now, with a new feature that will be the average of predictions by the bagging process. For the final model, I'll use cross validation and will perform feature reduction one more time, because we will have a new extra feature that will change the correlation between the data.

Joining our predictions into the datasets

```
In [56]:  
  
logreg = LogisticRegression(solver='newton-cg')  
param = {'C':[1,10]}  
LR = GridSearchCV(logreg,param,scoring='accuracy',refit=True,cv=10)  
LR.fit(x_train,y_train)  
score = LR.score(x_val, y_val)  
print(LR.best_params_,score)  
  
{'C': 1} 0.807495741056218
```

```
In [57]:  
  
x_train = pd.DataFrame(x_train)  
x_val = pd.DataFrame(x_val)
```

```
In [58]:  
  
x_train = x_train.join(train_pred)  
x_val = x_val.join(val_pred)
```

We will add more value joining those 3 columns by their mean

```
In [59]:  
  
x_train['proba'] = (x_train['LR'] + x_train['XGB'] + x_train['SVC']) / 3  
x_train = x_train.drop(['LR', 'XGB', 'SVC'], axis = 1)  
  
x_val['proba'] = (x_val['LR'] + x_val['XGB'] + x_val['SVC']) / 3  
x_val = x_val.drop(['LR', 'XGB', 'SVC'], axis = 1)
```

Let's check how does our model behaves with this new data

```
In [60]:  
  
logreg = LogisticRegression(solver='newton-cg')  
param = {'C':[1,10]}  
LR = GridSearchCV(logreg,param,scoring='accuracy',refit=True,cv=10)  
LR.fit(x_train,y_train)  
score = LR.score(x_val, y_val)  
print(LR.best_params_,score)  
  
{'C': 10} 0.817717206132879
```

Last step: one more feature reduction

```
In [61]:  
  
logisticRegr = LogisticRegression(solver='newton-cg')  
logisticRegr.fit(x_train, y_train)  
score = logisticRegr.score(x_val, y_val)  
print(score)  
  
pd.options.display.max_rows = 500  
  
importance = logisticRegr.coef_  
  
cols_importance = pd.DataFrame(importance)  
imp = cols_importance.T  
imp['columns'] = x_train.columns  
imp = imp.sort_values(by = 0, ascending = False).reset_index().drop(['index'], axis = 1)  
  
0.817717206132879
```

```
In [62]:
score_list_LR = []

for x in range(1,160):
    #Defining the columns set
    cols_2 = list(imp['columns'][0:x])
    x_train_LR = x_train[cols_2]
    x_val_LR = x_val[cols_2]
    #Score from LogReg
    logisticRegr = LogisticRegression(solver='newton-cg')
    logisticRegr.fit(x_train_LR, y_train)
    score = logisticRegr.score(x_val_LR, y_val)
    score_list_LR.append([score,x])
```

```
In [63]:
pd.options.display.max_rows = 50
top_LR = pd.DataFrame(score_list_LR).sort_values(0, ascending = False).head(5)
top_LR
```

	0	1
47	0.826235	48
17	0.824532	18
18	0.824532	19
33	0.822828	34
31	0.822828	32

```
In [64]:
cols_2 = list(imp['columns'][0:top_LR[1].iloc[0]])
x_train = x_train[cols_2]
x_val = x_val[cols_2]
```

Final model and score with CV and gridsearch:

```
In [65]:
logreg = LogisticRegression(solver='newton-cg')
param = {'C':[1,10]}
LR = GridSearchCV(logreg,param,scoring='accuracy',refit=True,cv=10)
LR.fit(x_train,y_train)
score = LR.score(x_val, y_val)
print(LR.best_params_,score)

{'C': 10} 0.8262350936967632
```

```
In [ ]:
```